

View Meta-Reviews

Paper ID

80

Paper Title

CODA: CASCADED ONLINE DISCONTINUITY-AWARE ALIGNMENT FOR REAL-TIME SCORE FOLLOWING

Track Name

Papers

META-REVIEWER #1

META-REVIEW QUESTIONS

2. I am an expert on the topic of the paper.

Agree

3. The title and abstract reflect the content of the paper.

Disagree

4. The paper discusses, cites and compares with all relevant related work.

Strongly agree

5. Please justify the previous choice (Required if “Strongly Disagree” or “Disagree” is chosen, otherwise write "n/a")

n/a

6. Readability and paper organization: The writing and language are clear and structured in a logical manner.

Strongly agree

7. The paper adheres to ISMIR 2026 submission guidelines (uses the ISMIR 2026 template, has at most 6 pages of technical content followed by “n” pages of references, AI usage declaration or ethical considerations, references are well formatted). If you selected “No”, please explain the issue in your comments.

Yes

8. Relevance of the topic to ISMIR: The topic of the paper is relevant to the ISMIR community. Note that submissions of novel music-related topics, tasks, and applications are highly encouraged. If you think that the paper has merit but does not exactly match the topics of ISMIR, please do not simply reject the paper but instead communicate this to the Program Committee Chairs. Please do not penalize the paper when the proposed method can also be applied to non-music domains if it is shown to be useful in music domains.

Strongly agree

9. Scholarly/scientific quality: The content is scientifically correct.

Strongly agree

10. Please justify the previous choice (Required if “Strongly Disagree” or “Disagree” is chose, otherwise write "n/a")

n/a

11. Novelty of the paper: The paper provides novel methods, applications, findings or results. Please do not narrowly view "novelty" as only new methods or theories. Papers proposing novel musical applications of existing methods from other research fields are considered novel at ISMIR conferences.

Strongly agree

12. The paper provides all the necessary details or material to reproduce the results described in the paper. Keep in mind that ISMIR respects the diversity of academic disciplines, backgrounds, and approaches. Although ISMIR has a tradition of publishing open datasets and open-source projects to enhance the scientific reproducibility, ISMIR accepts submissions using proprietary datasets and implementations that are not sharable. Please do not simply reject the paper when proprietary datasets or implementations are used.

Strongly agree

13. Pioneering proposals: This paper proposes a novel topic, task or application. Since this is intended to encourage brave new ideas and challenges, papers rated “Strongly Agree” and “Agree” can be highlighted, but please do not penalize papers rated “Disagree” or “Strongly Disagree”. Keep in mind that it is often difficult to provide baseline comparisons for novel topics, tasks, or applications. If you think that the novelty is high but the evaluation is weak, please do not simply reject the paper but carefully assess the value of the paper for the community.

Agree (Novel topic, task, or application)

14. Reusable insights: The paper provides reusable insights (i.e. the capacity to gain an accurate and deep understanding). Such insights may go beyond the scope of the paper, domain or application, in order to build up consistent knowledge across the MIR community.

Strongly agree

15. Please explain your assessment of reusable insights in the paper.

The cascaded score-to-audio alignment effectively addresses temporal discontinuity in music performance. The curriculum learning makes the cascaded model robust to potential failure from the top-level alignment.

16. AI Usage Policy: The paper complies with the ISMIR 2026 AI Usage Policy for Authors. (Authors are welcome to use any tool they wish for preparing and writing the paper but

they must ensure that all content is correct and original. Authors must declare any use of LLMs and AI tools if it's a part of core methodology (including literature review, related work, generated figures, tables or other illustrations). Using AI tools for editing purposes (e.g., checking grammar and fixing typos) need not be declared. Please pay attention to any hallucinated text, figures, or references and flag them in the comments below)

Agree

17. Please justify the previous choice (Required if "Disagree" is chosen, otherwise write "n/a")

n/a

18. Write ONE line (in your own words) with the main take-home message from the paper.

The cascaded score-to-audio alignment is robust to temporal discontinuity in music performance.

21. Potential to generate discourse: The paper will generate discourse at the ISMIR conference or have a large influence/impact on the future of the ISMIR community.

Agree

22. Overall evaluation (to be completed before the discussion phase): Please first evaluate before the discussion phase. Keep in mind that minor flaws can be corrected, and should not be a reason to reject a paper. Please familiarize yourself with the reviewer guidelines at <https://ismir.net/reviewer-guidelines>.

Strong accept

23. Main review and comments for the authors (to be completed before the discussion phase). Please summarize strengths and weaknesses of the paper. It is essential that you justify the reason for the overall evaluation score in detail. Keep in mind that belittling or sarcastic comments are not appropriate.

This paper presents a real-time score following method using a cascaded score-image-to-audio alignment based on the structure of music score. The proposed methods uses three-levels of alignments in system, bar, and notes and conducts a two-phase curriculum learning which prevents the model from cascaded failure from the upper level. The model is particularly robust in discontinuous jumps during the performance, which is frequent in practice scenarios. The experiment results shows outperforming alignment accuracy over compared models and the ablation study clarify the effectiveness of the model components.

This paper is very well written, and I do not find any significant issues that need to be addressed.

The three-level cascaded model is well-grounded based on practical score-and-audio alignment scenarios. The potential weakness of the model is remedied well by the two-phase curriculum training.

This paper has a very solid experimental setup that rigorously evaluates the proposed model with directly comparable models. The ablation study clearly shows the contribution of the model components. The dedicated jump recovery evaluation convincingly shows the advantage of the cascaded model.

A minor issue is that the title does not indicate that the score is an image modality from the title. Considering the majority of score following is based on audio-to-audio or midi-to-midi alignment, it would be good to specify that the score is represented as an image in the title.

Overall, this paper is well written and easy to read. The proposed model represents a significant step in score-image-to-audio alignment. It has a great potential to real-world score following in music education contexts. I strongly recommend acceptance of this paper.

24. Final recommendation (to be completed after the discussion phase) Please give a final recommendation after the discussion phase. In the final recommendation, please do not simply average the scores of the reviewers. Note that the number of recommendation options for reviewers is different from the number of options here. We encourage you to take a stand, and preferably avoid “weak accepts” or “weak rejects” if possible.

Strong accept

25. Meta-review and final comments for authors (to be completed after the discussion phase)

The reviewers unanimously praise the paper for its high-quality writing, strong musical grounding, and thorough evaluation.

The proposed two-phase curriculum learning strategy is highly commended for effectively mitigating the risk of cascaded failures from upper-level prediction errors. Furthermore, the model demonstrates state-of-the-art accuracy on the MSMD dataset and introduces a valuable repeat/jump-aware benchmark to the community, proving particularly robust against performance discontinuities.

There are a few minor points for clarification. They could be addressed in the camera-ready version.

Given the solid experimental validation, compelling ablation studies, and high potential for real-world application in music education, this paper is strongly recommended for acceptance. However, none of the reviews recommend it for a paper award.
